

dl-projectile-solver

January 9, 2020

0.1 Deep Learning Projectile Motion Solver

This notebook implements an artificial neural network to solve projectile motion problems when only three of the variables x , y , v_{0x} , v_{0y} , v_y , and t are known. A solution engine utilizing kinematic equations of motion is used to calculate exact solutions to problems used to both train the neural network and evaluate the accuracy of its results. The domain of the problems is constrained to simplify the coding of the exact solution engine. Specifically, v_{0x} and v_{0y} are always ≥ 0 , y is always ≤ 0 , and v_y is always an unknown.

0.1.1 Import useful packaged including TensorFlow 2.0 and Keras.

The notebook utilizes tensorflow ≥ 2.0 , which now includes keras, a package of high level wrappers designed to make building and training deep learning models easier.

```
[1]: # import packages
import tensorflow as tf
import tensorflow.keras as keras
import numpy as np
import random
import matplotlib.pyplot as plt
from matplotlib.lines import Line2D

from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.models import Sequential
from tensorflow.keras import optimizers
```

0.1.2 Check to see if a Tensorflow is installed with GPU support and if a GPU is available.

```
[2]: if not tf.test.is_gpu_available():
    print('No GPU found. Training will be slower.')
else:
    print('Default GPU {} found.'.format(tf.test.gpu_device_name()))
```

Default GPU /device:GPU:0 found.

0.1.3 Define helper functions.

```
[3]: # randomly select a set of three knowns and set all unknowns initially to zero.
def seed_inputs():

    # need at least one known for the x-dimension, and one from the y-
    ↪ dimension, but three in all.
    # to make imaginary and multiple solutions impossible, voy and v0x are
    ↪ always >=0,
    # y is always <= 0, and vy is always unknown.
    variables = ['x', 'y', 'v0x', 'v0y', 'vy', 't']
    knowns = random.sample(['v0x', 'x'],1) + random.sample(['v0y', 'y'],1)
    remaining = set(['x', 'y', 'v0x', 'v0y', 't']).difference(set(knowns))
    knowns = knowns + random.sample(remaining,1)
    vals = [round(random.uniform(0,1000),2), round(random.uniform(0,-500),2),
            round(random.uniform(50,100),2), round(random.uniform(50,100),2),
            round(random.uniform(0,-100),2), round(random.uniform(5,20),2)]
    ins = {variables[i]: vals[i] for i in range(len(variables))}

    for key in ins:
        if not key in knowns:
            ins[key] = 0

    return knowns, ins

# engine to solve projectile motion problem using kinematic equations of motion.
# this may not be the most efficient or elegant implementation but it works.
def solve_projectile(knowns, ins):
    outs = dict.copy(ins)
    g = -9.8

    if 'x' in knowns and 'v0x' in knowns: # find t
        outs['t'] = outs['x']/outs['v0x']
        knowns = knowns + ['t']

    if not 'y' in knowns:
        if 'v0y' in knowns and 'vy' in knowns: # if not T
            outs['t'] = (outs['vy'] - outs['v0y'])/g
        elif 'v0y' in knowns and 't' in knowns: # if not VY
            outs['vy'] = outs['v0y'] + g*outs['t']
        elif 'vy' in knowns and 't' in knowns: # if not VOY
            outs['v0y'] = outs['vy'] - g*outs['t']
        outs['y'] = outs['v0y']*outs['t'] + 0.5*g*outs['t']**2
    elif 'y' in knowns and 't' in knowns:
        outs['v0y'] = (outs['y'] - 0.5*g*outs['t']**2)/outs['t']
        outs['vy'] = outs['v0y'] + g*outs['t']
```

```

elif 'y' in knowns and 'vy' in knowns:
    outs['v0y'] = np.sqrt(outs['vy']**2 - 2*g*outs['y']) # always positive
    outs['t'] = (outs['vy']-outs['v0y'])/g
elif 'y' in knowns and 'v0y' in knowns:
    outs['vy'] = -np.sqrt(outs['v0y']**2 + 2*g*outs['y'])
    outs['t'] = (outs['vy']-outs['v0y'])/g
if 'x' in knowns and not 'v0x' in knowns:
    outs['v0x'] = outs['x']/outs['t']

if 'v0x' in knowns and not 'x' in knowns:
    outs['x'] = outs['v0x']*outs['t']

for key in outs:
    outs[key] = round(outs[key],2)

return outs

# generate N projectile motion problems and their solutions.
def generate_data(N):
    x = np.zeros((N,6))
    y = np.zeros((N,6))
    for ii in range(N):
        knowns, ins = seed_inputs()
        outs = solve_projectile(knowns, ins)
        x[ii,:] = list(ins.values())
        y[ii,:] = list(outs.values())

    return x, y

# function to test the solution engine so individual solutions can be checked
# for correctness.
def test_solution_engine():
    variables = np.array(['x', 'y', 'v0x', 'v0y', 'vy', 't'])
    knowns, ins = seed_inputs()
    outs = solve_projectile(knowns, ins)

    unknowns = list(set(variables).difference(set(knowns)))

    ins = list(ins.values())
    outs = list(outs.values())

    print('knowns:')
    for each in knowns:
        ii = list(np.where(variables==each))[0][0]
        print(each, '=', ins[ii])

```

```

print('\nsolved unknowns:')
for each in unknowns:
    ii = list(np.where(variables==each))[0][0]
    print(each, '=', outs[ii])

return

# function to calculate a projectiles trajectory from v0x, v0y, and x
def calc_trajectory(solution):
    theta = np.arctan(solution['v0y']/solution['v0x'])
    v0 = np.sqrt(solution['v0x']**2 + solution['v0y']**2)
    xrange = np.arange(0,solution['x'],0.1)
    yrange = xrange*np.tan(theta) - 9.8*np.multiply(xrange,xrange)/(2*v0**2*np.
    ↪cos(theta)**2)

    return xrange, yrange

```

0.1.4 Run the following cell to test the solution engine.

```
[4]: test_solution_engine()
```

```

knowns:
v0x = 99.08
y = -266.76
v0y = 76.62

solved unknowns:
t = 18.57
x = 1839.78
vy = -105.35

```

0.1.5 Generate training, test, and validation data sets.

```
[5]: x_train, y_train = generate_data(50000)
x_test, y_test = generate_data(5000)
x_validate, y_validate = generate_data(5000)
```

0.1.6 Define the ANN model graph.

```
[6]: # learning_rate is a hyperparameter
learning_rate = 0.0001
```

```

# four hidden layers - three fully connected and a dropout layer - and an
→ output layer
# linear activation is utilized on the final layer to get numerical results
→ commensurate with the type of
# problem we are considering.
model = Sequential()
model.add(Dense(128, activation='relu', input_dim=6))
model.add(Dense(256, activation='relu'))
model.add(Dropout(rate=0.05))
model.add(Dense(256, activation='relu'))
model.add(Dense(6, activation='linear'))

# the error function utilized is the sum of the mean squared error.
model.compile(loss='mean_squared_error',
              optimizer=keras.optimizers.Adam(lr=learning_rate),
              metrics=['accuracy'])

class AccuracyHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        self.acc = []
        self.val_acc = []

    def on_epoch_end(self, batch, logs={}):
        self.acc.append(logs.get('accuracy'))
        self.val_acc.append(logs.get('val_accuracy'))

history = AccuracyHistory()

```

0.1.7 Train the model.

```

[7]: # hyperparameters
batch_size = 64
epochs = 64

# train the model
model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_validate, y_validate),
          callbacks=[history])
score = model.evaluate(x_test, y_test, verbose=0)
print(score)

```

Train on 50000 samples, validate on 5000 samples

Epoch 1/64
50000/50000 [=====] - 2s 44us/sample - loss:
226825.0267 - accuracy: 0.9069 - val_loss: 156662.0082 - val_accuracy: 0.9374

Epoch 2/64
50000/50000 [=====] - 1s 30us/sample - loss:
205561.2255 - accuracy: 0.9245 - val_loss: 154005.6033 - val_accuracy: 0.9298

Epoch 3/64
50000/50000 [=====] - 2s 31us/sample - loss:
203286.3120 - accuracy: 0.9226 - val_loss: 151467.9689 - val_accuracy: 0.9310

Epoch 4/64
50000/50000 [=====] - 1s 29us/sample - loss:
201363.5741 - accuracy: 0.9252 - val_loss: 149307.5514 - val_accuracy: 0.9360

Epoch 5/64
50000/50000 [=====] - 1s 30us/sample - loss:
199268.3278 - accuracy: 0.9322 - val_loss: 146704.5826 - val_accuracy: 0.9442

Epoch 6/64
50000/50000 [=====] - 1s 30us/sample - loss:
197005.7727 - accuracy: 0.9422 - val_loss: 143848.7816 - val_accuracy: 0.9544

Epoch 7/64
50000/50000 [=====] - 1s 30us/sample - loss:
194245.8185 - accuracy: 0.9510 - val_loss: 141031.0822 - val_accuracy: 0.9584

Epoch 8/64
50000/50000 [=====] - 2s 31us/sample - loss:
192099.9537 - accuracy: 0.9558 - val_loss: 137642.2182 - val_accuracy: 0.9634

Epoch 9/64
50000/50000 [=====] - 2s 31us/sample - loss:
189560.1260 - accuracy: 0.9592 - val_loss: 135169.8088 - val_accuracy: 0.9682

Epoch 10/64
50000/50000 [=====] - 1s 30us/sample - loss:
186618.8667 - accuracy: 0.9618 - val_loss: 131401.5580 - val_accuracy: 0.9684

Epoch 11/64
50000/50000 [=====] - 2s 30us/sample - loss:
184006.7029 - accuracy: 0.9632 - val_loss: 127511.9567 - val_accuracy: 0.9688

Epoch 12/64
50000/50000 [=====] - 2s 30us/sample - loss:
180792.4118 - accuracy: 0.9634 - val_loss: 124000.7154 - val_accuracy: 0.9696

Epoch 13/64
50000/50000 [=====] - 2s 32us/sample - loss:
177997.3625 - accuracy: 0.9639 - val_loss: 119403.0051 - val_accuracy: 0.9716

Epoch 14/64
50000/50000 [=====] - 2s 31us/sample - loss:
174703.7560 - accuracy: 0.9647 - val_loss: 114992.7253 - val_accuracy: 0.9696

Epoch 15/64
50000/50000 [=====] - 1s 30us/sample - loss:
171292.8219 - accuracy: 0.9651 - val_loss: 110787.3508 - val_accuracy: 0.9718

Epoch 16/64
50000/50000 [=====] - 2s 30us/sample - loss:
167286.9927 - accuracy: 0.9659 - val_loss: 105729.2192 - val_accuracy: 0.9714

Epoch 17/64
50000/50000 [=====] - 2s 31us/sample - loss:
163514.1763 - accuracy: 0.9652 - val_loss: 102613.4849 - val_accuracy: 0.9702

Epoch 18/64
50000/50000 [=====] - 2s 33us/sample - loss:
160216.6553 - accuracy: 0.9642 - val_loss: 97338.6637 - val_accuracy: 0.9710

Epoch 19/64
50000/50000 [=====] - 2s 31us/sample - loss:
156521.5830 - accuracy: 0.9647 - val_loss: 93666.7120 - val_accuracy: 0.9716

Epoch 20/64
50000/50000 [=====] - 2s 32us/sample - loss:
154229.6766 - accuracy: 0.9654 - val_loss: 89676.0055 - val_accuracy: 0.9674

Epoch 21/64
50000/50000 [=====] - 2s 32us/sample - loss:
150622.7727 - accuracy: 0.9646 - val_loss: 89117.3173 - val_accuracy: 0.9728

Epoch 22/64
50000/50000 [=====] - 2s 31us/sample - loss:
147872.7199 - accuracy: 0.9663 - val_loss: 82350.9531 - val_accuracy: 0.9740

Epoch 23/64
50000/50000 [=====] - 2s 32us/sample - loss:
144762.6736 - accuracy: 0.9665 - val_loss: 77609.5376 - val_accuracy: 0.9690

Epoch 24/64
50000/50000 [=====] - 1s 30us/sample - loss:
140882.5599 - accuracy: 0.9658 - val_loss: 75497.1645 - val_accuracy: 0.9726

Epoch 25/64
50000/50000 [=====] - 1s 30us/sample - loss:
139670.0150 - accuracy: 0.9656 - val_loss: 71643.9692 - val_accuracy: 0.9734

Epoch 26/64
50000/50000 [=====] - 2s 34us/sample - loss:
134934.0436 - accuracy: 0.9664 - val_loss: 68505.9194 - val_accuracy: 0.9746

Epoch 27/64
50000/50000 [=====] - 2s 31us/sample - loss:
132604.8318 - accuracy: 0.9667 - val_loss: 66166.1641 - val_accuracy: 0.9774

Epoch 28/64
50000/50000 [=====] - 1s 30us/sample - loss:
128116.7638 - accuracy: 0.9673 - val_loss: 63429.5566 - val_accuracy: 0.9770

Epoch 29/64
50000/50000 [=====] - 1s 30us/sample - loss:
127712.3832 - accuracy: 0.9671 - val_loss: 58173.8889 - val_accuracy: 0.9774

Epoch 30/64
50000/50000 [=====] - 1s 30us/sample - loss:
124037.8837 - accuracy: 0.9672 - val_loss: 61002.8354 - val_accuracy: 0.9716

Epoch 31/64
50000/50000 [=====] - 2s 31us/sample - loss:
122256.5941 - accuracy: 0.9678 - val_loss: 55248.4380 - val_accuracy: 0.9746

Epoch 32/64
50000/50000 [=====] - 2s 33us/sample - loss:
116502.9261 - accuracy: 0.9687 - val_loss: 51723.7252 - val_accuracy: 0.9782

Epoch 33/64
50000/50000 [=====] - 2s 34us/sample - loss:
113723.5618 - accuracy: 0.9697 - val_loss: 55236.0293 - val_accuracy: 0.9786

Epoch 34/64
50000/50000 [=====] - 2s 30us/sample - loss:
116418.9519 - accuracy: 0.9699 - val_loss: 43481.1027 - val_accuracy: 0.9814

Epoch 35/64
50000/50000 [=====] - 2s 31us/sample - loss:
111896.7498 - accuracy: 0.9706 - val_loss: 44488.4195 - val_accuracy: 0.9788

Epoch 36/64
50000/50000 [=====] - 2s 34us/sample - loss:
110896.6180 - accuracy: 0.9711 - val_loss: 39506.6367 - val_accuracy: 0.9802

Epoch 37/64
50000/50000 [=====] - 2s 34us/sample - loss:
108678.1261 - accuracy: 0.9704 - val_loss: 37298.3340 - val_accuracy: 0.9746

Epoch 38/64
50000/50000 [=====] - 2s 33us/sample - loss:
103491.8128 - accuracy: 0.9717 - val_loss: 37589.3822 - val_accuracy: 0.9826

Epoch 39/64
50000/50000 [=====] - 2s 32us/sample - loss:
103134.0231 - accuracy: 0.9699 - val_loss: 32574.8370 - val_accuracy: 0.9828

Epoch 40/64
50000/50000 [=====] - 2s 32us/sample - loss:
101741.5785 - accuracy: 0.9709 - val_loss: 33616.2368 - val_accuracy: 0.9826

Epoch 41/64
50000/50000 [=====] - 2s 32us/sample - loss: 99322.1724
- accuracy: 0.9711 - val_loss: 29588.1351 - val_accuracy: 0.9810

Epoch 42/64
50000/50000 [=====] - 2s 32us/sample - loss: 97225.4852
- accuracy: 0.9703 - val_loss: 30238.4709 - val_accuracy: 0.9812

Epoch 43/64
50000/50000 [=====] - 2s 32us/sample - loss: 96541.6528
- accuracy: 0.9702 - val_loss: 28331.8182 - val_accuracy: 0.9812

Epoch 44/64
50000/50000 [=====] - 2s 32us/sample - loss: 93585.9041
- accuracy: 0.9723 - val_loss: 24742.7860 - val_accuracy: 0.9812

Epoch 45/64
50000/50000 [=====] - 2s 31us/sample - loss: 91557.1265
- accuracy: 0.9722 - val_loss: 23917.4806 - val_accuracy: 0.9810

Epoch 46/64
50000/50000 [=====] - 2s 32us/sample - loss: 86673.8677
- accuracy: 0.9705 - val_loss: 24981.9717 - val_accuracy: 0.9794

Epoch 47/64
50000/50000 [=====] - 2s 31us/sample - loss: 87267.1953
- accuracy: 0.9706 - val_loss: 25024.6665 - val_accuracy: 0.9794

Epoch 48/64
50000/50000 [=====] - 2s 32us/sample - loss: 84550.4746
- accuracy: 0.9705 - val_loss: 26502.9428 - val_accuracy: 0.9786

Epoch 49/64
50000/50000 [=====] - 2s 31us/sample - loss: 83723.3058
- accuracy: 0.9702 - val_loss: 18467.7277 - val_accuracy: 0.9804

Epoch 50/64
50000/50000 [=====] - 2s 32us/sample - loss: 80032.0324
- accuracy: 0.9705 - val_loss: 17431.4114 - val_accuracy: 0.9808

Epoch 51/64
50000/50000 [=====] - 2s 31us/sample - loss: 81359.5918
- accuracy: 0.9710 - val_loss: 19448.4440 - val_accuracy: 0.9804

Epoch 52/64
50000/50000 [=====] - 2s 31us/sample - loss: 81454.9648
- accuracy: 0.9712 - val_loss: 18295.1779 - val_accuracy: 0.9806

Epoch 53/64
50000/50000 [=====] - 2s 31us/sample - loss: 73807.0647
- accuracy: 0.9713 - val_loss: 16917.5558 - val_accuracy: 0.9796

Epoch 54/64
50000/50000 [=====] - 2s 31us/sample - loss: 81913.1900
- accuracy: 0.9710 - val_loss: 23103.9868 - val_accuracy: 0.9802

Epoch 55/64
50000/50000 [=====] - 2s 31us/sample - loss: 73979.3202
- accuracy: 0.9711 - val_loss: 17476.6455 - val_accuracy: 0.9800

Epoch 56/64
50000/50000 [=====] - 2s 31us/sample - loss: 76632.6042
- accuracy: 0.9721 - val_loss: 17569.1226 - val_accuracy: 0.9832

Epoch 57/64
50000/50000 [=====] - 2s 31us/sample - loss: 72044.8257
- accuracy: 0.9719 - val_loss: 15260.2575 - val_accuracy: 0.9820

Epoch 58/64
50000/50000 [=====] - 2s 31us/sample - loss: 72098.7199
- accuracy: 0.9729 - val_loss: 14909.3710 - val_accuracy: 0.9742

Epoch 59/64
50000/50000 [=====] - 2s 31us/sample - loss: 66594.7389
- accuracy: 0.9734 - val_loss: 26212.2037 - val_accuracy: 0.9878

Epoch 60/64
50000/50000 [=====] - 2s 32us/sample - loss: 69767.9273
- accuracy: 0.9736 - val_loss: 16098.3586 - val_accuracy: 0.9806

Epoch 61/64
50000/50000 [=====] - 2s 31us/sample - loss: 63826.5568
- accuracy: 0.9738 - val_loss: 16299.9440 - val_accuracy: 0.9832

Epoch 62/64
50000/50000 [=====] - 2s 31us/sample - loss: 66000.5024
- accuracy: 0.9746 - val_loss: 39307.8232 - val_accuracy: 0.9874

Epoch 63/64
50000/50000 [=====] - 2s 31us/sample - loss: 67496.6115
- accuracy: 0.9749 - val_loss: 16970.3528 - val_accuracy: 0.9844

Epoch 64/64
50000/50000 [=====] - 2s 31us/sample - loss: 66953.0000
- accuracy: 0.9740 - val_loss: 26566.4506 - val_accuracy: 0.9838


[9655.62845546875, 0.984]

0.1.8 Graph the accuracy of the model using the training data set versus the validation data set.

```
[8]: print('Test loss:', score[0])
      print('Test accuracy:', score[1])
      plt.plot(range(1, epochs+1), history.acc, range(1, epochs+1), history.val_acc)
      plt.legend(['Train', 'Validate'])
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.show()
```

Test loss: 9655.62845546875

Test accuracy: 0.984



0.1.9 Visualize solutions to the projectile motion problem predicted by the ANN

The solutions to the projectile motion problems predicted by the trained neural network are not exact solutions. In fact, the nature of the model can result in even the knowns fed into the network being adjusted. So, each prediction made by the ANN from a set of three knowns actually defines a variety of trajectories, which are visualized here.

```

[12]: variables = np.array(['x', 'y', 'v0x', 'v0y', 'vy', 't'])
knowns, ins = seed_inputs()
outs = solve_projectile(knowns, ins)
outs_prime = model.predict([list(ins.values())])
outs_prime = {variables[i]: round(outs_prime[0][i],2) for i in
    ↪range(len(variables))} # convert to dict
knowns_perms = [['v0x', 'v0y', 'x'], ['v0x', 'v0y', 'y'], ['v0x', 'v0y', 't'],
    ['v0x', 'y', 'x'], ['v0x', 'y', 't'],
    ['x', 'v0y', 'y'], ['x', 'v0y', 't'],
    ['x', 'y', 't']]
plt.figure(figsize=(10,6))

xr, yr = calc_trajectory(outs)
actual = plt.plot(xr, yr, 'k', linewidth=5)


for each in knowns_perms:
    solution = solve_projectile(each, outs_prime)
    xrange, yrange = calc_trajectory(solution)
    plt.plot(xrange,yrange,'r', linestyle='dashed')

plt.xlabel('x (m)', fontsize=12)
plt.ylabel('y (m)', fontsize=12)

legend_elements = [Line2D([0], [0], color='k', linewidth=5), Line2D([0], [0],
    ↪color='r', linestyle='dashed')]
labels = ['actual','predicted']

plt.legend(legend_elements, labels, fontsize=12)
plt.savefig('projectile.png')
plt.show()

```



[]: