

# dl-completely-inelastic-collision-rnn

January 9, 2020

## 0.1 Deep Learning Completely Inelastic Collision Solver

This notebook implements a recurrent neural network to solve for the post collision velocity of a two-mass system where the collision is completely inelastic and the second mass is initially at rest. The model is formulated using a recurrent neural network (RNN) and its input is represented as a text string giving the masses and the initial velocity of the first mass. Exact solutions calculated using conservation of momentum equations are used to train the RNN and evaluate the accuracy of predictions. This notebook is inspired by the `addition_rnn.py` example included with Keras.

### 0.1.1 Import useful packaged including TensorFlow 2.0 and Keras.

The notebook utilizes tensorflow  $\geq 2.0$ , which now includes keras, a package of high level wrappers designed to make building and training deep learning models easier.

```
[1]: # import packages
import tensorflow as tf
import tensorflow.keras as keras
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.lines import Line2D

from tensorflow.keras.layers import RNN, LSTM, TimeDistributed, RepeatVector,
↳Dense, LSTMCell, Dropout
# LSTM doesn't work well since cuDNN is compiled with certain restrictions.
# So, here I will create LSTM layers by wrapping LSTMCell in RNN
from tensorflow.keras.models import Sequential
from tensorflow.keras import optimizers
```

### 0.1.2 Check to see if a Tensorflow is installed with GPU support and if a GPU is available.

```
[2]: if not tf.test.is_gpu_available():
    print('No GPU found. Training will be slower.')
else:
    print('Default GPU {} found.'.format(tf.test.gpu_device_name()))
```

Default GPU /device:GPU:0 found.

### 0.1.3 Define a class to encode and decode between a selection of characters and one-hot integer representations.

```
[3]: class CharacterTable(object):
    """Given a set of characters:
    + Encode them to a one-hot integer representation
    + Decode the one-hot or integer representation to their character output
    + Decode a vector of probabilities to their character output
    """
    def __init__(self, chars):
        """Initialize character table.
        # Arguments
            chars: Characters that can appear in the input.
        """
        self.chars = sorted(set(chars))
        self.char_indices = dict((c, i) for i, c in enumerate(self.chars))
        self.indices_char = dict((i, c) for i, c in enumerate(self.chars))

    def encode(self, C, num_rows):
        """One-hot encode given string C.
        # Arguments
            C: string, to be encoded.
            num_rows: Number of rows in the returned one-hot encoding. This is
                used to keep the # of rows for each data the same.
        """
        x = np.zeros((num_rows, len(self.chars)))
        for i, c in enumerate(C):
            x[i, self.char_indices[c]] = 1
        return x

    def decode(self, x, calc_argmax=True):
        """Decode the given vector or 2D array to their character output.
        # Arguments
            x: A vector or a 2D array of probabilities or one-hot
            → representations;
                or a vector of character indices (used with
            → `calc_argmax=False`).
            calc_argmax: Whether to find the character index with maximum
                probability, defaults to `True`.
        """
        if calc_argmax:
            x = x.argmax(axis=-1)
        return ''.join(self.indices_char[x] for x in x)
```

```
class colors:
    ok = '\033[92m'
    fail = '\033[91m'
    close = '\033[0m'
```

#### 0.1.4 Generate input and output character sequences.

The input is in the form  $\{m\}v\{v\}$  where the values in the  $\{ \}$ s are randomly drawn one or two character sequences representing, from left to right, the mass of object one, the mass of object two, and the initial velocity of object one. Object two is initially at rest.

The output is a character string representing the post collision velocity of the system after a head on totally inelastic collision.

```
[4]: # Parameters for the model and dataset.
num_problems = 80000 # number of sequences in the training and validation sets
digits = 2 # maximum number of digits for mass and velocity in the input
    ↪ sequence
lenans = 4 # number of characters in the answer sequence

REVERSE=False

# Maximum length of input is 'int + int' (e.g., '345+678'). Maximum length of
# int is DIGITS.
maxlen = digits + 1 + digits + 1 + digits

# All the numbers, plus sign and space for padding.
chars = '0123456789mv. '
ctable = CharacterTable(chars)

questions = []
expected = []
seen = set()
print('Generating data...')
while len(questions) < num_problems:
    f = lambda: int(''.join(np.random.choice(list('123456789'))
        for i in np.arange(np.random.randint(1, digits + 1))))
    a, b, c = f(), f(), f()
    # Skip any questions we've already seen
    key = tuple(sorted((a, b, c)))
    if key in seen:
        continue
    seen.add(key)
    # Pad the data with spaces such that it is always MAXLEN.
    q = '{}m{}v{}'.format(a, b, c)
    query = q + ' ' * (maxlen - len(q))
    ans = a*c/(a+b)
```

```

    if ans < 10:
        r = lenans-2
    elif ans < 100:
        r = lenans-3
    ans = str(round(a*c/(a+b),r))
    # Answers can be of maximum size LENANS.
    ans += '0' * (lenans - len(ans))
    questions.append(query)
    expected.append(ans)
print('Total momentum questions:', len(questions))

print('Vectorization...')
x = np.zeros((len(questions), maxlen, len(chars)), dtype=np.bool)
y = np.zeros((len(questions), lenans, len(chars)), dtype=np.bool)
for i, sentence in enumerate(questions):
    x[i] = ctable.encode(sentence, maxlen)
for i, sentence in enumerate(expected):
    y[i] = ctable.encode(sentence, lenans)

# Shuffle (x, y) in unison as the later parts of x will almost all be larger
# digits.
indices = np.arange(len(y))
np.random.shuffle(indices)
x = x[indices]
y = y[indices]

# set apart 20% for validation and test data that we never train over.
split_at = len(x) - len(x) // 5
(x_train, x_val) = x[:split_at], x[split_at:]
(y_train, y_val) = y[:split_at], y[split_at:]

split_at = len(x_val) - len(x_val) // 2
(x_val, x_test) = x_val[:split_at], x_val[split_at:]
(y_val, y_test) = y_val[:split_at], y_val[split_at:]

print('Training Data:')
print(x_train.shape)
print(y_train.shape)

print('Validation Data:')
print(x_val.shape)
print(y_val.shape)

print('Test Data:')
print(x_test.shape)
print(y_test.shape)

```

```

Generating data...
Total momentum questions: 80000
Vectorization...
Training Data:
(64000, 8, 14)
(64000, 4, 14)
Validation Data:
(8000, 8, 14)
(8000, 4, 14)
Test Data:
(8000, 8, 14)
(8000, 4, 14)

```

```

[5]: # show an example of the input and output character strings
ii = np.random.randint(0, 1000)
print('input:', questions[ii], 'output:', expected[ii])

```

```

input: 1m9v93    output: 9.30

```

```

[6]: # hyperparameters
HIDDEN_SIZE = 128
BATCH_SIZE = 256
LAYERS = 2
learning_rate = 0.001

print('Build model...')
model = Sequential()
# "encode" the input sequence using an RNN, producing an output of HIDDEN_SIZE.
# note: in a situation where your input sequences have a variable length, use
# input_shape=(None, num_feature).
# model.add(LSTM(HIDDEN_SIZE, input_shape=(MAXLEN, len(chars))))
model.add(RNN(LSTMCell(HIDDEN_SIZE), input_shape=(maxlen, len(chars))))

# as the decoder RNN's input, repeatedly provide with the last output of
# RNN for each time step. Repeat 'lenans' times as that's the maximum length of
# output.
model.add(RepeatVector(lenans))
# the decoder RNN could be multiple layers stacked or a single layer.
for _ in range(LAYERS):
    # by setting return_sequences to True, return not only the last output but
    # all the outputs so far in the form of (num_samples, timesteps,
    # output_dim). This is necessary as TimeDistributed in the below expects
    # the first dimension to be the timesteps.
    # model.add(LSTM(HIDDEN_SIZE, return_sequences=True))
    model.add(RNN(LSTMCell(HIDDEN_SIZE), return_sequences=True))

# add a dropout layer to prevent overfitting

```

```

# model.add(Dropout(rate=0.1))
# add a dense layer to every temporal slice of an input. for each of step of
↳ the output sequence,
# decide which character should be chosen.
model.add(TimeDistributed(Dense(len(chars), activation='softmax'))))
model.compile(loss='categorical_crossentropy',
              # optimizer=keras.optimizers.Adam(lr=learning_rate),
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])
model.summary()

```

Build model...

Model: "sequential"

Layer (type)	Output Shape	Param #
rnn (RNN)	(None, 128)	73216
repeat_vector (RepeatVector)	(None, 4, 128)	0
rnn_1 (RNN)	(None, 4, 128)	131584
rnn_2 (RNN)	(None, 4, 128)	131584
time_distributed (TimeDistri	(None, 4, 14)	1806

Total params: 338,190  
 Trainable params: 338,190  
 Non-trainable params: 0

### 0.1.5 Train the model and generate example predictions every 3 epochs.

```

[7]: # train the model and show predictions against the validation dataset.
for iteration in range(1, 36):
    print()
    print('-' * 50)
    print('Iteration', iteration)
    model.fit(x_train, y_train,
              batch_size=BATCH_SIZE,
              epochs=3,
              validation_data=(x_val, y_val))
    score = model.evaluate(x_test, y_test, verbose=0)
    print('score:', score)
    # select 10 samples from the validation set at random so we can visualize
↳ errors.

```

```

for i in range(10):
    ind = np.random.randint(0, len(x_val))
    rowx, rowy = x_test[np.array([ind])], y_test[np.array([ind])]
    preds = model.predict_classes(1.0*rowx, verbose=0) # mult by 1.0 to get
    ↳data types to jive
    q = ctable.decode(rowx[0])
    correct = ctable.decode(rowy[0])
    guess = ctable.decode(preds[0], calc_argmax=False)
    print('Q', q[:-1] if REVERSE else q, end=' ')
    print('A', correct, end=' ')

    for ii in np.arange(len(correct)):
        if correct[ii] == guess[ii]:
            print(colors.ok + ' ' + colors.close, end='')
        else:
            print(colors.fail + ' ' + colors.close, end='')

    print(' P', guess)

```

-----  
Iteration 1

Train on 64000 samples, validate on 8000 samples

Epoch 1/3

64000/64000 [=====] - 6s 91us/sample - loss: 2.0252 -  
accuracy: 0.3221 - val\_loss: 1.9072 - val\_accuracy: 0.3399

Epoch 2/3

64000/64000 [=====] - 2s 39us/sample - loss: 1.8240 -  
accuracy: 0.3581 - val\_loss: 1.6968 - val\_accuracy: 0.3859

Epoch 3/3

64000/64000 [=====] - 2s 39us/sample - loss: 1.6521 -  
accuracy: 0.3887 - val\_loss: 1.5840 - val\_accuracy: 0.3912

score: [1.5855826959609984, 0.39]

Q 3m92v97	A 3.06	P 1.30
Q 94m53v88	A 56.3	P 44.6
Q 74m74v4	A 2.00	P 1.38
Q 29m32v31	A 14.7	P 11.3
Q 82m7v32	A 29.5	P 24.8
Q 59m98v4	A 1.50	P 0.33
Q 4m89v74	A 3.18	P 1.38
Q 77m16v75	A 62.1	P 54.6
Q 81m77v49	A 25.1	P 14.3
Q 41m9v74	A 60.7	P 58.8

-----  
Iteration 2

Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.5172 -  
accuracy: 0.4214 - val\_loss: 1.4270 - val\_accuracy: 0.4591  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 1.4187 -  
accuracy: 0.4543 - val\_loss: 1.4477 - val\_accuracy: 0.4371  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 1.3556 -  
accuracy: 0.4751 - val\_loss: 1.3870 - val\_accuracy: 0.4625  
score: [1.38504066324234, 0.46590626]  
Q 29m2v13 A 12.2 P 10.8  
Q 58m76v27 A 11.7 P 13.8  
Q 28m15v41 A 26.7 P 20.8  
Q 16m8v61 A 40.7 P 48.0  
Q 65m57v31 A 16.5 P 18.8  
Q 5m27v97 A 15.2 P 15.8  
Q 16m8v22 A 14.7 P 17.8  
Q 68m5v56 A 52.2 P 50.8  
Q 18m59v42 A 9.82 P 12.8  
Q 61m85v66 A 27.6 P 30.8

-----  
Iteration 3  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.3155 -  
accuracy: 0.4893 - val\_loss: 1.3143 - val\_accuracy: 0.4773  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.2835 -  
accuracy: 0.5013 - val\_loss: 1.2968 - val\_accuracy: 0.4835  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.2592 -  
accuracy: 0.5097 - val\_loss: 1.2515 - val\_accuracy: 0.5061  
score: [1.25640402507782, 0.5069375]  
Q 52m76v4 A 1.62 P 1.80  
Q 26m31v82 A 37.4 P 44.2  
Q 23m25v14 A 6.71 P 6.80  
Q 43m34v8 A 4.47 P 4.20  
Q 78m48v1 A 0.62 P 0.62  
Q 92m1v15 A 14.8 P 16.8  
Q 89m86v44 A 22.4 P 24.2  
Q 5m57v6 A 0.48 P 0.42  
Q 4m81v7 A 0.33 P 0.42  
Q 24m24v23 A 11.5 P 12.2

-----  
Iteration 4  
Train on 64000 samples, validate on 8000 samples



Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.2372 -  
accuracy: 0.5167 - val\_loss: 1.2266 - val\_accuracy: 0.5182  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.2172 -  
accuracy: 0.5231 - val\_loss: 1.2430 - val\_accuracy: 0.5013  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1993 -  
accuracy: 0.5298 - val\_loss: 1.2064 - val\_accuracy: 0.5277  
score: [1.2089817209243774, 0.526875]  
Q 63m58v25 A 13.0 P 12.2  
Q 37m58v34 A 13.2 P 13.2  
Q 45m75v36 A 13.5 P 13.2  
Q 47m91v45 A 15.3 P 14.2  
Q 45m4v23 A 21.1 P 21.6  
Q 84m54v21 A 12.8 P 12.2  
Q 62m42v81 A 48.3 P 49.6  
Q 81m41v23 A 15.3 P 14.2  
Q 9m57v27 A 3.68 P 3.70  
Q 55m68v87 A 38.9 P 35.7

-----  
Iteration 5  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1831 -  
accuracy: 0.5356 - val\_loss: 1.1539 - val\_accuracy: 0.5506  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1664 -  
accuracy: 0.5411 - val\_loss: 1.1241 - val\_accuracy: 0.5673  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1503 -  
accuracy: 0.5472 - val\_loss: 1.1327 - val\_accuracy: 0.5453  
score: [1.1363491096496583, 0.54075]  
Q 38m85v21 A 6.49 P 6.11  
Q 78m59v97 A 55.2 P 56.5  
Q 99m5v67 A 63.8 P 61.1  
Q 21m31v8 A 3.23 P 3.50  
Q 33m87v16 A 4.40 P 4.80  
Q 16m71v22 A 4.05 P 4.11  
Q 16m8v22 A 14.7 P 13.4  
Q 89m3v82 A 79.3 P 78.1  
Q 5m85v12 A 0.67 P 0.75  
Q 88m24v29 A 22.8 P 21.8

-----  
Iteration 6  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 41us/sample - loss: 1.1355 -  
accuracy: 0.5518 - val\_loss: 1.1521 - val\_accuracy: 0.5374  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1223 -  
accuracy: 0.5569 - val\_loss: 1.0885 - val\_accuracy: 0.5785  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.1066 -  
accuracy: 0.5634 - val\_loss: 1.1267 - val\_accuracy: 0.5455  
score: [1.1247176299095154, 0.545625]  
Q 82m19v44 A 35.7 P 37.3  
Q 41m13v78 A 59.2 P 61.0  
Q 31m35v31 A 14.6 P 15.3  
Q 39m54v21 A 8.81 P 8.10  
Q 75m58v85 A 47.9 P 50.3  
Q 21m8v18 A 13.0 P 14.8  
Q 75m3v92 A 88.5 P 89.0  
Q 83m1v13 A 12.8 P 13.8  
Q 84m73v32 A 17.1 P 17.3  
Q 48m7v58 A 50.6 P 51.8

-----  
Iteration 7  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.0948 -  
accuracy: 0.5667 - val\_loss: 1.0560 - val\_accuracy: 0.5927  
Epoch 2/3  
64000/64000 [=====] - 3s 41us/sample - loss: 1.0814 -  
accuracy: 0.5710 - val\_loss: 1.0823 - val\_accuracy: 0.5720  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 1.0692 -  
accuracy: 0.5775 - val\_loss: 1.0657 - val\_accuracy: 0.5742  
score: [1.0661763501167298, 0.5701875]  
Q 14m39v26 A 6.87 P 6.60  
Q 38m58v28 A 11.1 P 11.1  
Q 92m97v29 A 14.1 P 13.1  
Q 59m36v4 A 2.48 P 2.10  
Q 98m76v38 A 21.4 P 20.1  
Q 1m11v45 A 3.75 P 3.70  
Q 8m23v35 A 9.03 P 7.00  
Q 14m89v97 A 13.2 P 13.1  
Q 22m54v78 A 22.6 P 22.3  
Q 98m88v24 A 12.6 P 12.1

-----  
Iteration 8  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0580 -  
accuracy: 0.5813 - val\_loss: 1.1565 - val\_accuracy: 0.5343  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0475 -  
accuracy: 0.5848 - val\_loss: 1.0371 - val\_accuracy: 0.5832  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0366 -  
accuracy: 0.5898 - val\_loss: 1.0094 - val\_accuracy: 0.5998  
score: [1.0074513502120972, 0.6014063]  
Q 99m23v56 A 45.4 P 44.5  
Q 61m2v94 A 91.0 P 91.1  
Q 46m97v44 A 14.2 P 13.9  
Q 64m89v62 A 25.9 P 25.5  
Q 91m1v27 A 26.7 P 26.6  
Q 8m18v43 A 13.2 P 12.0  
Q 38m2v71 A 67.5 P 68.9  
Q 22m8v73 A 53.5 P 54.5  
Q 1m51v22 A 0.42 P 0.25  
Q 41m34v63 A 34.4 P 33.5

-----  
Iteration 9  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0279 -  
accuracy: 0.5941 - val\_loss: 1.0148 - val\_accuracy: 0.5924  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0197 -  
accuracy: 0.5962 - val\_loss: 1.0253 - val\_accuracy: 0.5934  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 1.0099 -  
accuracy: 0.6014 - val\_loss: 1.0736 - val\_accuracy: 0.5699  
score: [1.0733536467552185, 0.568375]  
Q 2m25v31 A 2.30 P 2.90  
Q 41m76v58 A 20.3 P 20.8  
Q 41m13v78 A 59.2 P 69.9  
Q 69m66v62 A 31.7 P 32.9  
Q 8m48v61 A 8.71 P 8.00  
Q 67m6v62 A 56.9 P 58.1  
Q 24m72v15 A 3.75 P 3.12  
Q 61m7v13 A 11.7 P 12.4  
Q 94m43v89 A 61.1 P 61.9  
Q 41m59v11 A 4.51 P 4.12

-----  
Iteration 10  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 1.0022 -  
accuracy: 0.6037 - val\_loss: 0.9463 - val\_accuracy: 0.6365  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.9947 -  
accuracy: 0.6069 - val\_loss: 0.9755 - val\_accuracy: 0.6210  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.9860 -  
accuracy: 0.6104 - val\_loss: 0.9564 - val\_accuracy: 0.6224  
score: [0.9550179929733277, 0.6228125]  
Q 41m34v63 A 34.4 P 33.5  
Q 24m48v81 A 27.0 P 26.6  
Q 36m37v7 A 3.45 P 3.60  
Q 96m64v86 A 51.6 P 50.6  
Q 99m18v29 A 24.5 P 23.6  
Q 15m75v73 A 12.2 P 12.6  
Q 13m75v62 A 9.16 P 9..2  
Q 7m89v7 A 0.51 P 0.42  
Q 91m26v1 A 0.78 P 0.76  
Q 49m98v9 A 3.00 P 2.56

-----  
Iteration 11  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.9791 -  
accuracy: 0.6124 - val\_loss: 0.9770 - val\_accuracy: 0.6109  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.9716 -  
accuracy: 0.6165 - val\_loss: 1.0375 - val\_accuracy: 0.5902  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.9655 -  
accuracy: 0.6185 - val\_loss: 0.9217 - val\_accuracy: 0.6428  
score: [0.9180031900405884, 0.6445625]  
Q 48m17v79 A 58.3 P 58.8  
Q 84m15v47 A 39.9 P 49.8  
Q 13m37v59 A 15.3 P 15.1  
Q 61m59v12 A 6.10 P 6.10  
Q 68m16v33 A 26.7 P 26.1  
Q 58m13v95 A 77.6 P 76.1  
Q 39m97v91 A 26.1 P 25.1  
Q 9m17v49 A 17.0 P 16.8  
Q 77m49v3 A 1.83 P 1.78  
Q 56m44v28 A 15.7 P 15.1

-----  
Iteration 12  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.9587 -  
accuracy: 0.6226 - val\_loss: 0.9403 - val\_accuracy: 0.6226  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.9515 -  
accuracy: 0.6248 - val\_loss: 1.0730 - val\_accuracy: 0.5728  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.9478 -  
accuracy: 0.6260 - val\_loss: 1.0166 - val\_accuracy: 0.6011  
score: [1.0162253673076629, 0.601375]  
Q 86m99v32 A 14.9 P 15.8  
Q 89m86v44 A 22.4 P 22.8  
Q 35m6v39 A 33.3 P 33.8  
Q 85m51v63 A 39.4 P 38.0  
Q 46m65v77 A 31.9 P 32.8  
Q 83m7v45 A 41.5 P 41.8  
Q 41m2v99 A 94.4 P 94.8  
Q 26m52v85 A 28.3 P 28.2  
Q 16m57v4 A 0.88 P 0.94  
Q 73m19v16 A 12.7 P 12.8

-----  
Iteration 13  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.9416 -  
accuracy: 0.6291 - val\_loss: 0.9205 - val\_accuracy: 0.6395  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.9360 -  
accuracy: 0.6313 - val\_loss: 0.8947 - val\_accuracy: 0.6570  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.9322 -  
accuracy: 0.6318 - val\_loss: 0.8813 - val\_accuracy: 0.6644  
score: [0.8812202425003052, 0.66353124]  
Q 93m8v11 A 10.1 P 10.1  
Q 73m28v2 A 1.45 P 1.43  
Q 24m83v43 A 9.64 P 9..0  
Q 13m75v56 A 8.27 P 8.90  
Q 13m91v78 A 9.75 P 9..1  
Q 52m9v1 A 0.85 P 0.80  
Q 33m6v64 A 54.2 P 54.0  
Q 58m55v63 A 32.3 P 32.0  
Q 57m93v74 A 28.1 P 28.0  
Q 77m18v43 A 34.9 P 35.0

-----  
Iteration 14  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.9266 -  
accuracy: 0.6358 - val\_loss: 0.9004 - val\_accuracy: 0.6448  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.9201 -  
accuracy: 0.6375 - val\_loss: 0.9676 - val\_accuracy: 0.6114  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.9159 -  
accuracy: 0.6390 - val\_loss: 0.9426 - val\_accuracy: 0.6218  
score: [0.9460811560153961, 0.61928123]  
Q 75m6v58 A 53.7 P 53.2  
Q 31m58v62 A 21.6 P 21.2  
Q 3m72v15 A 0.60 P 0.62  
Q 97m65v73 A 43.7 P 42.2  
Q 25m74v55 A 13.9 P 13.2  
Q 41m2v99 A 94.4 P 94.0  
Q 24m42v92 A 33.5 P 33.2  
Q 12m79v69 A 9.10 P 8.22  
Q 8m21v72 A 19.9 P 29.2  
Q 23m11v85 A 57.5 P 56.2

-----  
Iteration 15  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.9127 -  
accuracy: 0.6403 - val\_loss: 0.9325 - val\_accuracy: 0.6245  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.9073 -  
accuracy: 0.6424 - val\_loss: 0.8621 - val\_accuracy: 0.6674  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.9047 -  
accuracy: 0.6437 - val\_loss: 0.8590 - val\_accuracy: 0.6737  
score: [0.8556039447784424, 0.6741875]  
Q 56m49v91 A 48.5 P 48.0  
Q 62m75v51 A 23.1 P 23.6  
Q 4m98v18 A 0.71 P 0.72  
Q 3m81v14 A 0.50 P 0.53  
Q 61m91v2 A 0.80 P 0.86  
Q 22m88v1 A 0.20 P 0.22  
Q 8m33v85 A 16.6 P 15.2  
Q 8m71v15 A 1.52 P 1.53  
Q 61m53v6 A 3.21 P 3.10  
Q 92m25v38 A 29.9 P 29.6

-----  
Iteration 16  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8982 -  
accuracy: 0.6468 - val\_loss: 0.9359 - val\_accuracy: 0.6243  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8949 -  
accuracy: 0.6485 - val\_loss: 0.8677 - val\_accuracy: 0.6587  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8915 -  
accuracy: 0.6490 - val\_loss: 0.9218 - val\_accuracy: 0.6384  
score: [0.919591515302658, 0.6401875]  
Q 6m91v58 A 3.59 P 3.00  
Q 9m9v55 A 27.5 P 27.0  
Q 48m56v72 A 33.2 P 32.0  
Q 27m32v94 A 43.0 P 42.0  
Q 49m74v24 A 9.56 P 9.40  
Q 71m37v57 A 37.5 P 37.0  
Q 17m57v29 A 6.66 P 6.00  
Q 19m62v33 A 7.74 P 7.60  
Q 11m45v22 A 4.32 P 4.70  
Q 28m91v23 A 5.41 P 5.02

-----  
Iteration 17  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8881 -  
accuracy: 0.6504 - val\_loss: 0.8694 - val\_accuracy: 0.6616  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8826 -  
accuracy: 0.6533 - val\_loss: 0.8683 - val\_accuracy: 0.6633  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8805 -  
accuracy: 0.6540 - val\_loss: 0.8797 - val\_accuracy: 0.6520  
score: [0.879545749425888, 0.6511875]  
Q 26m42v64 A 24.5 P 24.5  
Q 7m67v5 A 0.47 P 0.48  
Q 52m82v24 A 9.31 P 9.60  
Q 94m47v36 A 24.0 P 23.8  
Q 8m37v43 A 7.64 P 7.50  
Q 71m89v99 A 43.9 P 44.5  
Q 53m59v28 A 13.2 P 13.5  
Q 15m42v2 A 0.53 P 0.40  
Q 43m14v25 A 18.9 P 19.5  
Q 7m87v98 A 7.30 P 6.50

-----  
Iteration 18  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8775 -  
accuracy: 0.6547 - val\_loss: 0.8412 - val\_accuracy: 0.6785  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.8729 -  
accuracy: 0.6578 - val\_loss: 0.9125 - val\_accuracy: 0.6407  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8697 -  
accuracy: 0.6592 - val\_loss: 0.9355 - val\_accuracy: 0.6315  
score: [0.9333925821781158, 0.631375]  
Q 32m8v9 A 7.20 P 6.58  
Q 29m26v86 A 45.3 P 44.9  
Q 5m24v24 A 4.14 P 3.55  
Q 6m12v8 A 2.67 P 2.80  
Q 96m95v8 A 4.02 P 4.90  
Q 35m33v65 A 33.5 P 33.5  
Q 98m77v6 A 3.36 P 3.51  
Q 53m75v32 A 13.2 P 13.5  
Q 35m5v4 A 3.50 P 3.47  
Q 72m52v41 A 23.8 P 24.8

-----  
Iteration 19  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8688 -  
accuracy: 0.6578 - val\_loss: 0.8574 - val\_accuracy: 0.6658  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8644 -  
accuracy: 0.6605 - val\_loss: 0.8219 - val\_accuracy: 0.6862  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8597 -  
accuracy: 0.6615 - val\_loss: 0.8497 - val\_accuracy: 0.6683  
score: [0.8487990510463714, 0.66809374]  
Q 81m48v35 A 22.0 P 22.8  
Q 75m96v11 A 4.82 P 4.88  
Q 28m75v33 A 8.97 P 9.10  
Q 2m65v67 A 2.00 P 2.08  
Q 75m6v29 A 26.9 P 26.8  
Q 16m83v99 A 16.0 P 16.8  
Q 19m42v23 A 7.16 P 7.80  
Q 9m11v32 A 14.4 P 14.0  
Q 97m23v61 A 49.3 P 48.8  
Q 7m89v94 A 6.85 P 6.88

-----  
Iteration 20  
Train on 64000 samples, validate on 8000 samples



Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8561 -  
accuracy: 0.6644 - val\_loss: 0.9543 - val\_accuracy: 0.6270  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.8556 -  
accuracy: 0.6645 - val\_loss: 0.8521 - val\_accuracy: 0.6663  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8514 -  
accuracy: 0.6655 - val\_loss: 0.8926 - val\_accuracy: 0.6453  
score: [0.8878784575462342, 0.64821875]  
Q 57m72v64 A 28.3 P 28.5  
Q 34m92v85 A 22.9 P 23.1  
Q 85m51v63 A 39.4 P 39.5  
Q 26m54v5 A 1.62 P 1.65  
Q 92m2v39 A 38.2 P 38.1  
Q 9m7v46 A 25.9 P 26.5  
Q 65m3v68 A 65.0 P 65.2  
Q 76m8v51 A 46.1 P 46.5  
Q 74m36v83 A 55.8 P 55.1  
Q 36m59v22 A 8.34 P 8.60

-----  
Iteration 21  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8488 -  
accuracy: 0.6674 - val\_loss: 0.8952 - val\_accuracy: 0.6420  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8433 -  
accuracy: 0.6698 - val\_loss: 0.8314 - val\_accuracy: 0.6712  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8416 -  
accuracy: 0.6701 - val\_loss: 0.8651 - val\_accuracy: 0.6618  
score: [0.866441558599472, 0.6590313]  
Q 22m54v57 A 16.5 P 16.0  
Q 8m35v13 A 2.42 P 2.20  
Q 41m48v8 A 3.69 P 3.50  
Q 52m74v12 A 4.95 P 4.30  
Q 78m46v29 A 18.2 P 18.8  
Q 2m89v29 A 0.64 P 0.53  
Q 28m41v73 A 29.6 P 29.3  
Q 75m3v92 A 88.5 P 88.5  
Q 37m81v79 A 24.8 P 24.3  
Q 71m14v5 A 4.18 P 4.18

-----  
Iteration 22  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8392 -  
accuracy: 0.6709 - val\_loss: 0.8508 - val\_accuracy: 0.6631  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8361 -  
accuracy: 0.6733 - val\_loss: 0.8439 - val\_accuracy: 0.6679  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8347 -  
accuracy: 0.6732 - val\_loss: 0.8815 - val\_accuracy: 0.6482  
score: [0.8830624108314514, 0.6481562]  
Q 59m57v24 A 12.2 P 11.0  
Q 53m6v8 A 7.19 P 7.25  
Q 7m23v17 A 3.97 P 3.00  
Q 47m18v98 A 70.9 P 71.0  
Q 28m88v8 A 1.93 P 1.80  
Q 98m91v18 A 9.33 P 9.10  
Q 55m97v86 A 31.1 P 30.3  
Q 24m9v17 A 12.4 P 12.1  
Q 57m75v51 A 22.0 P 21.1  
Q 43m21v7 A 4.70 P 4.10

-----  
Iteration 23  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.8305 -  
accuracy: 0.6741 - val\_loss: 0.8411 - val\_accuracy: 0.6668  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8275 -  
accuracy: 0.6763 - val\_loss: 0.8227 - val\_accuracy: 0.6740  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8267 -  
accuracy: 0.6771 - val\_loss: 0.8481 - val\_accuracy: 0.6644  
score: [0.8486135015487671, 0.6670625]  
Q 2m53v65 A 2.36 P 2.20  
Q 76m84v66 A 31.4 P 31.8  
Q 22m13v41 A 25.8 P 25.0  
Q 17m3v45 A 38.2 P 38.8  
Q 97m94v66 A 33.5 P 33.0  
Q 34m38v27 A 12.8 P 12.8  
Q 6m8v59 A 25.3 P 24.0  
Q 97m44v29 A 20.0 P 20.8  
Q 5m95v11 A 0.55 P 0.43  
Q 18m38v61 A 19.6 P 18.0

-----  
Iteration 24  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8223 -  
accuracy: 0.6785 - val\_loss: 0.8180 - val\_accuracy: 0.6837  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8195 -  
accuracy: 0.6793 - val\_loss: 0.8330 - val\_accuracy: 0.6690  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8179 -  
accuracy: 0.6807 - val\_loss: 0.7950 - val\_accuracy: 0.6927  
score: [0.7929418609142304, 0.6969375]  
Q 47m67v49 A 20.2 P 20.5  
Q 76m8v51 A 46.1 P 46.8  
Q 46m69v93 A 37.2 P 37.0  
Q 38m51v35 A 14.9 P 15.6  
Q 38m18v19 A 12.9 P 12.8  
Q 2m22v77 A 6.42 P 6.00  
Q 14m22v42 A 16.3 P 16.0  
Q 81m66v21 A 11.6 P 11.5  
Q 68m6v54 A 49.6 P 59.8  
Q 47m9v66 A 55.4 P 55.5

-----  
Iteration 25  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8155 -  
accuracy: 0.6811 - val\_loss: 0.7967 - val\_accuracy: 0.6908  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8133 -  
accuracy: 0.6815 - val\_loss: 0.8581 - val\_accuracy: 0.6598  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8105 -  
accuracy: 0.6831 - val\_loss: 0.8080 - val\_accuracy: 0.6808  
score: [0.8066634085178376, 0.68290627]  
Q 95m54v58 A 37.0 P 37.5  
Q 95m16v35 A 30.0 P 30.8  
Q 5m24v63 A 10.9 P 11.5  
Q 81m8v16 A 14.6 P 14.6  
Q 4m37v53 A 5.17 P 5.00  
Q 61m34v52 A 33.4 P 33.5  
Q 71m85v51 A 23.2 P 23.5  
Q 92m8v78 A 71.8 P 71.5  
Q 12m53v9 A 1.66 P 1.55  
Q 94m86v39 A 20.4 P 20.5

-----  
Iteration 26  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.8074 -  
accuracy: 0.6841 - val\_loss: 0.8086 - val\_accuracy: 0.6776  
Epoch 2/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.8052 -  
accuracy: 0.6858 - val\_loss: 0.8475 - val\_accuracy: 0.6694  
Epoch 3/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.8019 -  
accuracy: 0.6864 - val\_loss: 0.8037 - val\_accuracy: 0.6816  
score: [0.8009607315063476, 0.6852813]  
Q 67m44v71 A 42.9 P 43.0  
Q 96m37v56 A 40.4 P 40.8  
Q 67m32v67 A 45.3 P 45.0  
Q 47m98v7 A 2.27 P 2.33  
Q 72m2v4 A 3.89 P 3.89  
Q 57m23v52 A 37.0 P 37.0  
Q 62m94v54 A 21.5 P 21.5  
Q 6m76v76 A 5.56 P 5.50  
Q 84m98v34 A 15.7 P 16.1  
Q 12m96v29 A 3.22 P 3.30

-----  
Iteration 27  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7995 -  
accuracy: 0.6881 - val\_loss: 0.8097 - val\_accuracy: 0.6853  
Epoch 2/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.7986 -  
accuracy: 0.6888 - val\_loss: 0.8176 - val\_accuracy: 0.6776  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7955 -  
accuracy: 0.6890 - val\_loss: 0.8158 - val\_accuracy: 0.6735  
score: [0.8188097815513611, 0.67175]  
Q 26m5v34 A 28.5 P 28.8  
Q 62m26v86 A 60.6 P 60.9  
Q 23m15v7 A 4.24 P 4.20  
Q 98m4v41 A 39.4 P 39.4  
Q 5m68v84 A 5.75 P 5.09  
Q 37m27v56 A 32.4 P 32.9  
Q 83m12v43 A 37.6 P 36.9  
Q 24m72v59 A 14.8 P 15.2  
Q 67m91v44 A 18.7 P 18.9  
Q 53m59v28 A 13.2 P 13.6

-----  
Iteration 28  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7933 -  
accuracy: 0.6895 - val\_loss: 0.8693 - val\_accuracy: 0.6583  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7918 -  
accuracy: 0.6907 - val\_loss: 0.7983 - val\_accuracy: 0.6832  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7881 -  
accuracy: 0.6927 - val\_loss: 0.8026 - val\_accuracy: 0.6790  
score: [0.8030582339763641, 0.6786875]  
Q 25m29v65 A 30.1 P 30.6  
Q 66m24v5 A 3.67 P 3.66  
Q 51m96v93 A 32.3 P 32.0  
Q 68m81v19 A 8.67 P 8.60  
Q 7m22v43 A 10.4 P 10.6  
Q 25m53v75 A 24.0 P 23.3  
Q 35m35v8 A 4.00 P 3.90  
Q 58m59v98 A 48.6 P 49.5  
Q 37m47v85 A 37.4 P 37.0  
Q 62m15v3 A 2.42 P 2.40

-----  
Iteration 29  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.7872 -  
accuracy: 0.6931 - val\_loss: 0.7769 - val\_accuracy: 0.6988  
Epoch 2/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7840 -  
accuracy: 0.6938 - val\_loss: 0.8692 - val\_accuracy: 0.6630  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7817 -  
accuracy: 0.6953 - val\_loss: 0.8320 - val\_accuracy: 0.6703  
score: [0.8350286138057709, 0.66809374]  
Q 9m47v78 A 12.5 P 12.8  
Q 16m19v1 A 0.46 P 0.46  
Q 72m27v99 A 72.0 P 71.6  
Q 34m38v27 A 12.8 P 12.8  
Q 4m56v28 A 1.87 P 1.90  
Q 54m58v39 A 18.8 P 18.5  
Q 66m58v31 A 16.5 P 16.6  
Q 48m6v2 A 1.78 P 1.75  
Q 56m12v71 A 58.5 P 57.6  
Q 94m84v61 A 32.2 P 31.0

-----  
Iteration 30  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.7793 -  
accuracy: 0.6959 - val\_loss: 0.7784 - val\_accuracy: 0.6968  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7776 -  
accuracy: 0.6969 - val\_loss: 0.7880 - val\_accuracy: 0.6925  
Epoch 3/3  
64000/64000 [=====] - 2s 38us/sample - loss: 0.7744 -  
accuracy: 0.6982 - val\_loss: 0.8171 - val\_accuracy: 0.6817  
score: [0.8186441586017609, 0.67934376]  
Q 84m2v74 A 72.3 P 72.2  
Q 47m98v35 A 11.3 P 11.0  
Q 28m9v36 A 27.2 P 27.0  
Q 3m45v68 A 4.25 P 4.25  
Q 7m67v5 A 0.47 P 0.44  
Q 18m95v79 A 12.6 P 12.0  
Q 53m78v82 A 33.2 P 32.5  
Q 41m57v7 A 2.93 P 2.90  
Q 2m33v3 A 0.17 P 0.15  
Q 13m43v83 A 19.3 P 19.2

-----  
Iteration 31  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 3s 40us/sample - loss: 0.7739 -  
accuracy: 0.6985 - val\_loss: 0.7901 - val\_accuracy: 0.6886  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7712 -  
accuracy: 0.6992 - val\_loss: 0.7590 - val\_accuracy: 0.7062  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.7711 -  
accuracy: 0.6995 - val\_loss: 0.7859 - val\_accuracy: 0.6930  
score: [0.7823098633289337, 0.698]  
Q 56m7v86 A 76.4 P 76.1  
Q 79m57v92 A 53.4 P 53.3  
Q 52m32v25 A 15.5 P 15.5  
Q 42m14v87 A 65.2 P 65.5  
Q 4m92v28 A 1.17 P 1.15  
Q 58m89v35 A 13.8 P 13.8  
Q 81m88v63 A 30.2 P 30.3  
Q 2m61v3 A 0.10 P 0.09  
Q 24m79v78 A 18.2 P 18.5  
Q 15m42v2 A 0.53 P 0.50

-----  
Iteration 32  
Train on 64000 samples, validate on 8000 samples

Epoch 1/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7674 -  
accuracy: 0.7008 - val\_loss: 0.8750 - val\_accuracy: 0.6531  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7644 -  
accuracy: 0.7033 - val\_loss: 0.8079 - val\_accuracy: 0.6768  
Epoch 3/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7633 -  
accuracy: 0.7032 - val\_loss: 0.7945 - val\_accuracy: 0.6870  
score: [0.7964618723392487, 0.68703127]  
Q 38m9v29 A 23.4 P 23.0  
Q 93m5v58 A 55.0 P 55.1  
Q 6m81v47 A 3.24 P 3.20  
Q 24m12v82 A 54.7 P 55.9  
Q 79m68v7 A 3.76 P 3.77  
Q 16m83v9 A 1.45 P 1.45  
Q 99m45v5 A 3.44 P 3.51  
Q 93m36v46 A 33.2 P 33.2  
Q 21m8v18 A 13.0 P 12.5  
Q 26m38v18 A 7.31 P 7.80

-----  
Iteration 33  
Train on 64000 samples, validate on 8000 samples  
Epoch 1/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.7609 -  
accuracy: 0.7042 - val\_loss: 0.8159 - val\_accuracy: 0.6761  
Epoch 2/3  
64000/64000 [=====] - 3s 39us/sample - loss: 0.7588 -  
accuracy: 0.7043 - val\_loss: 0.7773 - val\_accuracy: 0.6978  
Epoch 3/3  
64000/64000 [=====] - 2s 39us/sample - loss: 0.7567 -  
accuracy: 0.7058 - val\_loss: 0.8803 - val\_accuracy: 0.6568  
score: [0.8798099870681763, 0.65578127]  
Q 67m23v24 A 17.9 P 17.8  
Q 37m61v22 A 8.31 P 8.46  
Q 3m83v7 A 0.24 P 0.24  
Q 43m69v27 A 10.4 P 10.1  
Q 34m79v98 A 29.5 P 29.2  
Q 67m87v1 A 0.44 P 0.45  
Q 47m44v21 A 10.8 P 10.8  
Q 61m8v39 A 34.5 P 35.6  
Q 35m33v28 A 14.4 P 14.0  
Q 9m26v3 A 0.77 P 0.72

-----  
Iteration 34  
Train on 64000 samples, validate on 8000 samples

```

Epoch 1/3
64000/64000 [=====] - 3s 40us/sample - loss: 0.7554 -
accuracy: 0.7063 - val_loss: 0.7768 - val_accuracy: 0.6930
Epoch 2/3
64000/64000 [=====] - 3s 40us/sample - loss: 0.7540 -
accuracy: 0.7064 - val_loss: 0.7572 - val_accuracy: 0.7030
Epoch 3/3
64000/64000 [=====] - ETA: 0s - loss: 0.7515 -
accuracy: 0.70 - 2s 38us/sample - loss: 0.7521 - accuracy: 0.7061 - val_loss:
0.7697 - val_accuracy: 0.6945
score: [0.7694476416110992, 0.69578123]
Q 31m14v27 A 18.6    P 18.5
Q 57m26v17 A 11.7    P 11.6
Q 7m25v34  A 7.44     P 7.20
Q 25m6v91  A 73.4     P 74.8
Q 88m35v38 A 27.2     P 27.0
Q 47m2v98  A 94.0     P 94.2
Q 96m7v62  A 57.8     P 58.1
Q 88m7v91  A 84.3     P 84.3
Q 19m57v36 A 9.00     P 8.12
Q 68m17v3  A 2.40     P 2.41

```

```

-----
Iteration 35
Train on 64000 samples, validate on 8000 samples
Epoch 1/3
64000/64000 [=====] - 2s 39us/sample - loss: 0.7482 -
accuracy: 0.7089 - val_loss: 0.7972 - val_accuracy: 0.6853
Epoch 2/3
64000/64000 [=====] - 2s 38us/sample - loss: 0.7480 -
accuracy: 0.7082 - val_loss: 0.7576 - val_accuracy: 0.7023
Epoch 3/3
64000/64000 [=====] - 2s 39us/sample - loss: 0.7467 -
accuracy: 0.7091 - val_loss: 0.7545 - val_accuracy: 0.7038
score: [0.7542344851493835, 0.7033125]
Q 98m94v95 A 48.5    P 48.0
Q 36m62v35 A 12.9    P 12.6
Q 44m28v8  A 4.89     P 4.70
Q 55m8v53  A 46.3     P 46.5
Q 3m93v4   A 0.12     P 0.12
Q 9m48v83  A 13.1     P 12.4
Q 65m26v67 A 47.9     P 48.0
Q 92m72v59 A 33.1     P 32.0
Q 61m56v73 A 38.1     P 37.0
Q 9m58v42  A 5.64     P 5.40

```

[ ]: