

W20: Deep Learning with Python

Saturday January 12, 2019, 1:00 - 5:00 pm

AAPT Winter Meeting, Houston, Texas


Jeff Groff, Shepherd University

Workshop Materials:

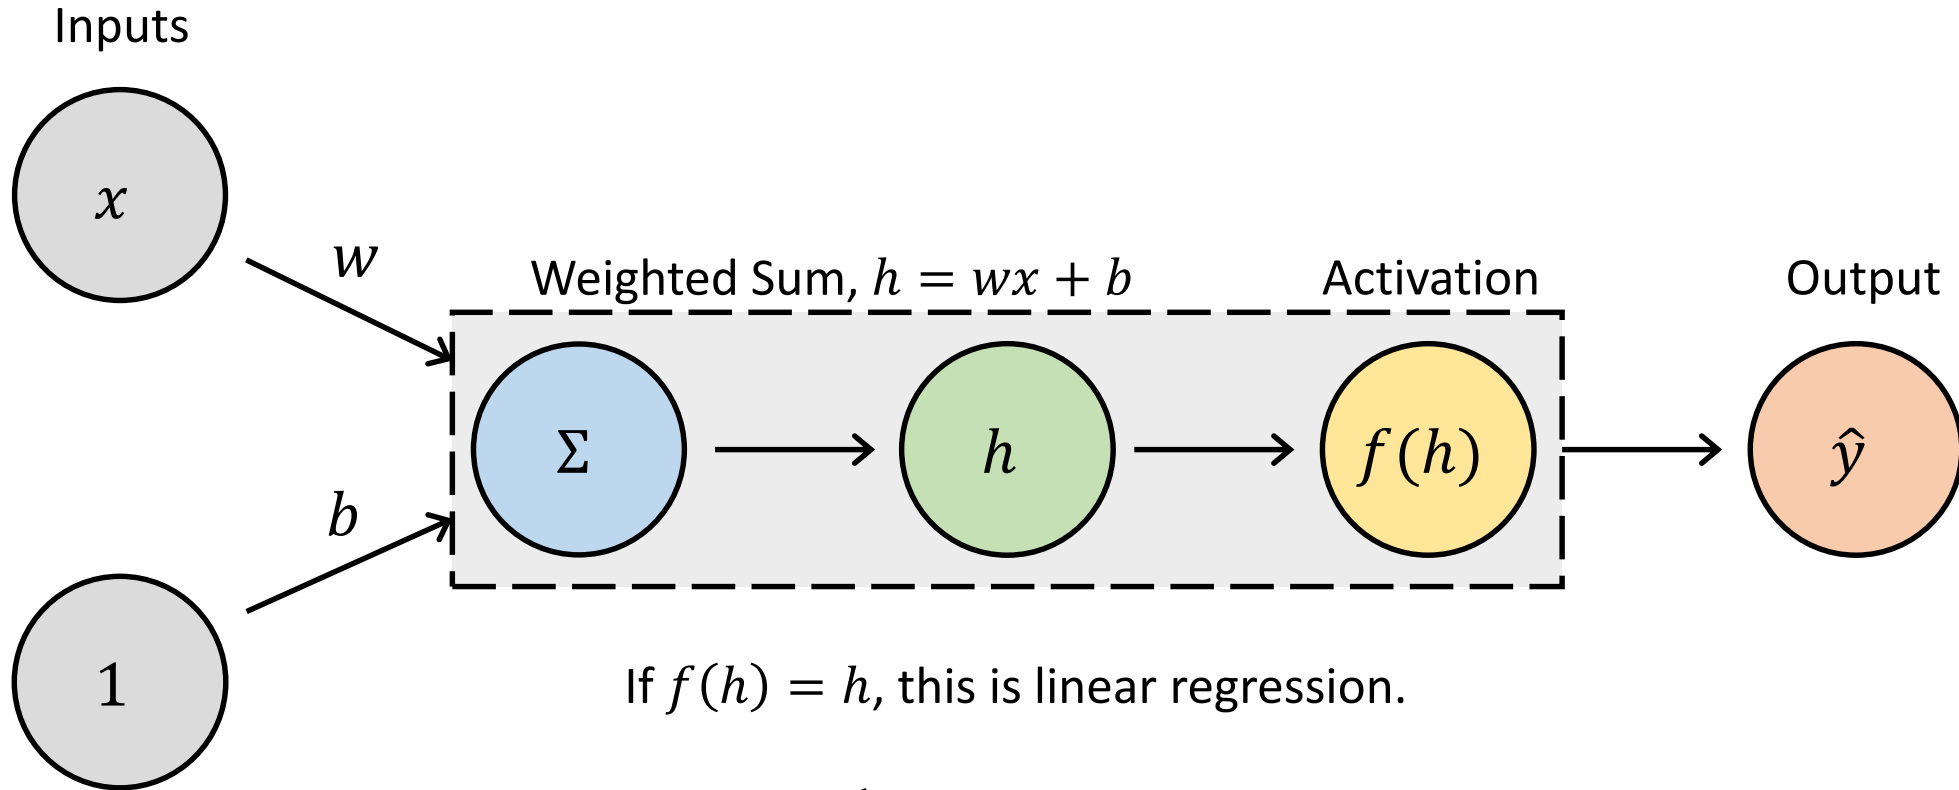
<https://github.com/prof-groff/deep-learning>



Deep Learning: A Venn Diagram



Single-Layer, Single-Node Neural Network



If $f(h) = h$, this is linear regression.

If $f(h) = \begin{cases} 1 & \text{if } h > 0 \\ 0 & \text{otherwise} \end{cases}$, this is a perceptron.

If $f(h) = \frac{1}{1 + e^{-h}}$, this is logistic regression.


Gradient Descent

Goal: Find network weights that minimize the error between the model outputs (predictions) and the actual data (targets).

To minimize $E(w, b) = \frac{1}{2} (y - \hat{y}(w, b))^2$ we can iterate each weight at each time step by a value proportional to the partial derivative of the error with respect to the weight.

$$\Delta w = w_{n+1} - w_n = -\eta \frac{\partial E}{\partial w}$$

$$\Delta b = b_{n+1} - b_n = -\eta \frac{\partial E}{\partial b}$$




$$\Delta w = \frac{-\eta \sum_{\text{all data}} \frac{\partial E}{\partial w}}{\text{number of data records}}$$

Gradient Descent

$$E = \frac{1}{2}(y - \hat{y})^2 \quad \frac{\partial E}{\partial w} = -(y - \hat{y}) \frac{\partial \hat{y}}{\partial w} = -(y - \hat{y}) f'(h) \frac{\partial h}{\partial w} = -(y - \hat{y}) f'(h) x$$

$$\frac{\partial E}{\partial b} = -(y - \hat{y}) f'(h)$$




$f'(h) = 1$, for linear regression.

$f'(h) = f(h)(1 - f(h))$, for logistic regression.

Multi-Node, Multi-Layer Neural Network

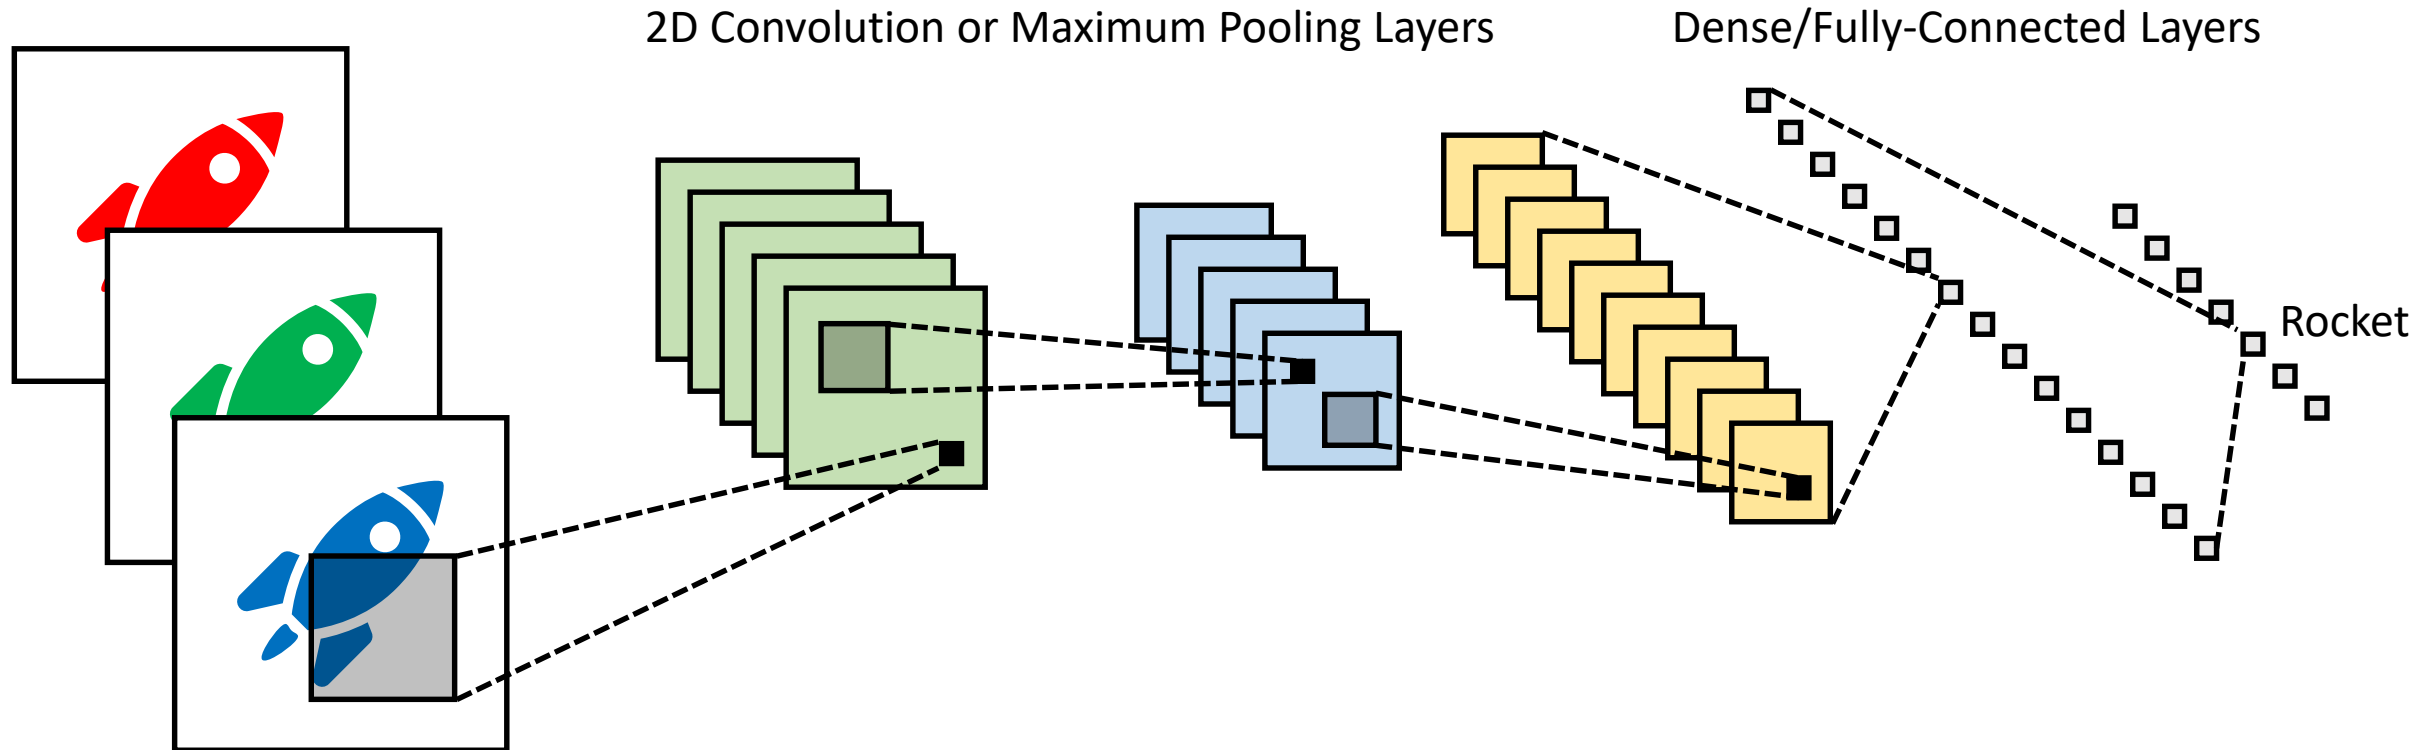
Deep Learning Models Have Many Layers




Convolutional Neural Network

Input Layer:
3-Channel
(RGB) Image

Output Layer:
One Node for
Each Class



Recurrent Neural Network



Unrolled